

Sona — Product Strategy & Differentiation Review

May 2026

2026-05-28

Contents

TL;DR — answers in 60 seconds

Context

1. Where the one-therapist tailoring leaks through

2. Unique selling point and moat

One-sentence USP

One-paragraph positioning

Comparison table — what an SLT actually evaluates

Moat analysis

3. The seams to add so Sona generalises without forking

Seam 1 — Intake template engine (S → M)

Seam 2 — Plan and summary template library (M)

Seam 3 — Vocabulary and statutory-framework pack (S)

Seam 4 — Multi-jurisdiction plumbing extensions (M)

Seam 5 — Multi-tenant clinic mode (M → L)

Seam 6 — White-label / per-practice branding (S)

Seam 7 — Specialty-aware persona library (S, ongoing)

4. External-systems integration roadmap

UK (priority)

US (Year 2+)

AU / CA (Year 2+, brief)

Phased summary

5. Pricing options

Cost-side reality check

Pricing options scored

Recommended starting model

Recommendations — what to do in the next 30 / 60 / 90 days

Open questions for the founders

Sources

Internal source artefacts (non-public, team repository)

Web sources (URL + access date, all 2026-05-28)

Engineering footnotes

TL;DR — answers in 60 seconds

- **Single-customer fit risk is real and concentrated.** The 8-step parent form, the four triage outcomes, and every persona are paediatric, UK, English-speaking, EHCP-aware, and Speech-Sanctuary-branded. A second private SLT — especially adult voice, AAC, dysphagia, or stutter — will spot the tailoring inside one screen. **Fixable by configuration, not by re-engineering**, but the configuration layer does not yet exist.
- **USP in one line.** *Sona turns the first 30 days of a private speech-therapy case — intake, triage, first-session plan, parent summary — into one reviewed loop, with UK-region data and a clinician sign-off on every artifact.* The defensible part is the **loop** and the **review-gated trust posture**, not any single AI feature.
- **Moat ranking, honest.** Strongest: design-partner-shaped intake-to-summary loop and the audit-trail-by-construction posture. Medium: UK data residency and the structured, schema-validated AI output shapes. **Weakest: the self-hosted Gemma posture itself** — competitors will match it within twelve months, and most clinicians cannot tell a self-hosted Gemma from a BAA-covered API.
- **Integration headline.** Adjacent integrations (calendar, video, e-sign, payments, and a "Sona for Cliniko" wedge) are the realistic Year-1 path. **NHS GP Connect is a Year-2+ programme**, not a Year-1 feature — it requires DSPT Standards Met, HSCN access, PDS compliance, a Clinical Safety Officer compliant with DCB0129/0160, SCAL evidence, and use-case approval. (NHS Digital, *GP Connect Access Record: Structured*, retrieved 2026-05-28.)
- **Pricing recommendation.** Per-clinician monthly subscription with a free-trial-on-cases mechanic. **Start at £79 / clinician / month** including AI drafting on a usage cap, with a £39 "starter" tier for new graduates. Sits between Smilenotes (£5, no AI) and Heidi Clinician (\$150 ≈ £120, scribe-only). Per-case is wrong for this market: 3–8 cases / week / clinician puts revenue too close to a single Cliniko seat to be defensible.
- **What to do next week.** (1) Run the one-therapist-fit audit with a second SLT before any code lands. (2) Decide whether the next sprint is the intake template engine or the AI loop, and only one — both at once is the failure mode. (3) Price a DSPT readiness pass with a specialist consultancy. (4) Talk to ASLTIP about a member-benefit slot.

Context

The MVP brief and the demo deck pitch Sona as a co-pilot for the first 30 days of a private SLT-to-family relationship: smart parent intake → consult prep → triage → first-session plan → parent summary. The demo video and slides — the artifacts a reviewer will have seen — quote a **30–60 minute saving per case at 3–8 cases per week**, which lines up with the per-stage time-saving targets the team has set internally for the triage step, the session-plan step, and the parent-summary step.

The MVP was scoped from a 1.5-hour interview with **one** private SLT (the design partner). That is enough to ship a demo; it is not enough to know whether the next ten therapists who watch the same demo see themselves on the screen. That is the question this memo exists to answer.

What is **real today**: the 8-step parent intake form, case management, audit trail, the design system, and slot-based booking. What is **vision**: AI prep brief, AI plan, AI parent summary, telehealth, calendar sync. Anything below that relies on the vision side is flagged as such.

1. Where the one-therapist tailoring leaks through

The audit walks the intake spec, the AI output shapes, the triage outcomes, and the vocabulary choices. Each surface is classified by **who it fits**, **who it breaks**, and a **fix class** — *config* (a per-practice setting toggle in existing infrastructure), *template* (a new template family), or *branch* (a fork in the product).

Surface	Assumption today	Fits	Breaks for	Fix class
Intake form: 8 hard-coded steps in the Flutter app	One global question tree	Paediatric speech-sounds and language clients	Adult voice, adult stutter, AAC, dysphagia, transgender voice — each needs its own tree	template engine
Form header copy: "Confidential Speech & Language Parent/Carer Questionnaire", tagline "Breathe / Speak / Be"	Speech Sanctuary brand	Speech Sanctuary	Every other practice — visibly someone else's form	config (white-label)
Mother / Father labelled fields, page 1	Two-parent UK paediatric household	Most paediatric cases	Single-parent, same-sex parents, kinship carers, looked-after children, all adult clients	config (rename → "Parent / carer 1 / 2") + template variants
GP details as free text (practice, address, phone)	UK	UK paediatric	US (PCP), AU (GP via My Health Record), CA (family physician) — and even UK if/when GP Connect provides structured data	config + integration (see roadmap)
EHCP yes/no + EHCP-aware prompt path	English statutory framework	England-based school-age clients	Wales (IDP), Scotland (CSP), NI (Statement), Republic of Ireland, US (IEP under IDEA), AU (NDIS) — different statutory wrappers entirely	template (statutory-framework pack)
Nursery / SENCO referral language	Early-years UK	Aria-shaped paediatric cases	Adult clients (post-stroke, head & neck cancer, voice, gender-affirming voice)	config (referral-source enum) + template (adult intake)
Difficulty checklist of 21 items across 4 domains (speech & language, attention, social, reading & writing)	Paediatric general SLT	Paediatric general SLT	AAC users (no symbol-set / device questions); feeding clients (no food-texture / mealtime log); transgender voice (no goals framing for adult voice transition); stutter sub-specialty (no severity rating scale)	template per specialty
Triage outcomes: a 4-way enum (strategy only, short block, full assessment, refer out)	The design partner's private-practice funnel	Solo UK private practice	NHS-funded routes; US insurance-coded plans; AU NDIS pathway — different funnels entirely	config (per-tenant outcome enum)

Surface	Assumption today	Fits	Breaks for	Fix class
Session-plan output schema's specialty enum: speech sounds, language, social communication, fluency, voice, feeding, other	Broad paediatric SLT	Most paediatric and some adult	Dysphagia/swallow as a first-class domain; AAC; tracheostomy / laryngectomy voice; transgender voice goals	template (extend enum + per-domain prompt fragments)
Parent-summary tone slider (warm ↔ clinical, 5 points) + reading-level (primary/secondary/adult)	Parent of a child	Almost everyone — generalises well	Pure adult-client cases where there is no parent ("client summary", not "parent summary")	config (rename to "recipient summary")
"Parent" used in field labels, audit events, UI strings	Paediatric model of care	UK paediatric private	Schools-based SLT (TA, SENCO), residential care (key worker), adult clinic (client / patient), public-sector NHS contexts (carer)	config (vocabulary pack)
Magic-link parent identity (no NHS Login)	Self-pay private	Self-pay private UK	Anyone using NHS-funded routes; ICB-commissioned services; school SLA workflows	integration (see roadmap)
Solo-practitioner default (from the MVP brief)	One therapist owns a tenant	Sole practitioners ≈ majority of ASLTIP's 1,800 members	Small group practices (the natural growth path for design partners); schools-based teams; charity providers	branch (multi-tenant clinic mode)
Regulators baked into copy: HCPC, RCSLT	UK	UK	US (ASHA), AU (SPA), CA (SAC), IE (CORU)	config (regulator pack tied to jurisdiction)
Demo personas: 3y feeding → 11y social-comm, all UK English	Paediatric demo coverage	Paediatric general SLT	Anything outside paediatric general SLT — the design partner has not been tested on adult, AAC, or English-as-additional-language cases	persona library expansion before any external pitch beyond paediatric

Three leaks matter more than the others, and each maps to a single concrete fix.

Leak 1: the form is the product, and the form is one practice's form. The 8-step intake is hard-coded in the Flutter app, and the API just persists whatever the form sends as an opaque blob — there is no form-template subsystem on either side. A second SLT cannot author her own form without an engineer. **Fix: an intake template engine** (seam 1 below). Until then, the strongest demo recommendation is to white-label the existing form for the *next* SLT — change the practice name, drop the Speech Sanctuary tagline — but **do not change the questions**, so the underlying assumption is visibly stress-tested.

Leak 2: paediatric is doing all the heavy lifting. The four personas, the EHCP wiring, the "parent" vocabulary, and the triage outcomes are all paediatric-shaped. Sona will demo well to a paediatric private SLT. It will demo badly to an adult voice SLT, an AAC specialist, or anyone running a feeding clinic. There is no shame in narrowing the pitch — **paediatric private SLT** is a credible wedge — but the pitch deck and the website have to say that out loud, not pretend Sona is for "speech and language therapy, less admin" in general.

Leak 3: the AI output schemas are tight, which is good for safety and bad for portability. The session-plan output schema enumerates seven specialty domains. Extending to dysphagia or AAC means schema + prompt + few-shot examples + UI badge. That is doable, but the *cost* per new specialty is non-trivial, and Sona should not invite an AAC clinician to a feedback session until that work is scoped.

2. Unique selling point and moat

One-sentence USP

Sona is the AI co-pilot for private speech-and-language therapy that takes the first 30 days of every new case — intake, triage, first-session plan, parent summary — from sixty minutes of admin per case to fifteen, with a clinician review on every output and UK data residency by default.

One-paragraph positioning

Private speech-and-language therapy is a 1,800-clinician UK profession in which most practitioners juggle a generic practice-management tool, a Word template, and ChatGPT-in-a-browser. The result: hours per week of admin, real client data leaving the UK, and no audit trail. Sona collapses the first 30 days of every case — adaptive parent intake, consult prep, four-outcome triage, first-session plan, parent summary — into one loop, with every AI artifact labelled and clinician-reviewed before it leaves the practice. The data and the inference both stay in the UK region. The output is structured by specialty and by statutory framework, not free-text-and-pray. The clinician owns every word that reaches the family.

Comparison table — what an SLT actually evaluates

The alternatives below are what a paediatric private SLT in 2026 is already weighing. Prices are headline, billed annually where annual pricing exists, **per practitioner per month** unless noted. Currencies preserved from the source.

Tool	What it does for an SLT	Headline price	Where it wins	Where it loses for the intake → first-session loop	Why Sona is different
ChatGPT / Claude / Gemini + a custom GPT	Free-text drafting; clinician copy-pastes from intake email	£0–£20/mo	Familiarity; zero adoption cost; very fast	PHI leaves the UK and the practice's control; no audit trail; no schema, every output drifts; no structured intake	UK-region inference; review-gated outputs; structured intake feeds the drafting context
Heidi Health	Ambient AI scribe — generates clinical notes from a session recording	\$150 / user / mo (Clinician, annual); BAA at Enterprise	Strong note quality across specialties; multilingual; mature	Scribes a session that has already happened — does nothing for intake, triage, or first-session drafting; no UK-specific data-residency claim by default	Sona owns the pre-session and post-session admin loop, which Heidi by design does not

Tool	What it does for an SLT	Headline price	Where it wins	Where it loses for the intake → first-session loop	Why Sona is different
Jane App AI Scribe	Recording-based scribe inside Jane PMS, SLP sample prompts	\$15 / mo unlimited , plus 5/mo free	Cheapest credible AI scribe; SLP-aware prompts; HIPAA/PIPEDA/PHIPA	Not available in the UK as of 2026-05-28; tied to Jane PMS; scribe-only	Available in the UK by construction; not a scribe — owns the structured first-30-days workflow
Cliniko + AI-scribe add-on (PatientNotes, CliniScripts, Clindoc, Heidi, Clinic Notes AI)	Generic PMS with bring-your-own scribe via marketplace	Cliniko from ~£35 + scribe \$19–\$150	The market default; calendar + invoicing + records already exist	Two products, two logins, two audit trails; intake is upload-a-PDF; no SLT-specific structure	Sona is one product with one audit trail and SLT-specific structure — and a "Sona for Cliniko" wedge is a wedge, not a competitor (see the integration roadmap)
Clindoc	UK-friendly AI scribe + clinical documentation, GDPR + HIPAA	£40 / user / mo annual or £5 / active day	UK GBP billing; per-active-day pricing for variable caseloads; lighter procurement	Scribe-only; no intake; no triage; no parent-facing artefacts	Sona is the workflow product the scribe-only tools complement
PatientNotes	AI clinical notes + integrations with Cliniko and Power Diary; SLT-aware marketing	\$19 / \$49 / mo	Cheap; honest about its scope; SLT-aware copy	Note-taker only; no parent-facing summary; no UK-only inference posture	Sona is end-to-end, not bring-your-own
Octopus EPR (OEPR)	UK SLT-specific EPR — casenotes, TOMs scores, RISK, programmes, discharge letters; ASLTIP-marketed	Pricing not posted	The most clinically aware UK SLT-specific competitor ; designed by SLTs; speaks the domain language	Records-keeping tool — no AI drafting; intake is not the focus; tools-not-loops	Sona owns the <i>new-case</i> loop; OEPR owns the <i>ongoing-record</i> loop. Different problem, today. Same RFP, eventually.
Smilenotes UK	UK SLT PMS; clinical notes templates; UK-hosted	From £5 / mo	Cheapest credible UK option; GDPR-compliant; UK data centres	No AI; basic templates; no triage / plan / parent-summary structure	Sona is the AI loop that sits on top of (or alongside) a tool like Smilenotes

Tool	What it does for an SLT	Headline price	Where it wins	Where it loses for the intake → first-session loop	Why Sona is different
SLT-specific AI entrants (Constant Therapy, Better Speech "Jessica", Ogma Therapy, IEP Copilot, SPRY)	A grab-bag of client-facing apps, IEP drafters, and US PT/SLP rehab platforms	Variable	Some are UK-friendly; some are clinically rich	Mostly client-facing kid-facing apps (Sona has an explicit principle against adding to children's screen time) or US-only EMR	Sona is adult-facing tooling that helps the SLT help the family, not another kid-app
The status quo (ChatGPT + Cliniko + Word template + email)	What every private SLT already does	Already paid for	Already adopted; already familiar	Hours of admin per week; PHI in ChatGPT; no audit trail	The thing Sona has to beat. The competitor in every meeting.
Doing nothing	—	£0	No change-management cost	Admin burnout is a documented driver of attrition out of private practice (per the design-partner interview)	Sona aims to be a measurable hour-a-week back

Moat analysis

No single item below is a 5-year moat on its own. The strength is in the **stack**. Time-to-erosion = how long until a well-funded competitor could match this if they prioritised it.

#	Candidate moat	Defensibility	Time-to-erosion	Read
1	Design-partner-shaped intake → triage → plan → summary loop	Medium-high	18–24 months	The end-to-end loop, structured by schema-validated output formats and clinician-review gates at each step, is the actual product. Replicating it requires a design-partner relationship <i>and</i> the discipline not to ship a generic scribe. Most competitors will ship the scribe.
2	Audit trail and review-gated AI as a posture, not a feature	Medium-high	18 months	An append-only audit log, clinician-review gates on every artifact, and a hard-coded AI-disclosure footer mean nothing can leave the system without a clinician act. That posture sells trust in a profession trained to fear AI. Heidi can match the controls; the <i>posture</i> — small clinician-first product, not enterprise scribe — is harder to copy.
3	Specialty-aware structured outputs	Medium	12–18 months	The seven-domain session-plan output, the EHCP flag, the four triage outcomes — these are specificity that a generic scribe will not invest in until the SLT segment is bigger. Becomes weaker as more competitors target SLT specifically (Jane App is already there for note quality).

#	Candidate moat	Defensibility	Time-to-erosion	Read
4	UK-region data residency + self-hosted Gemma in a London Google Cloud region	Medium-low	12 months	A genuine differentiator vs ChatGPT and US-only Jane App. But — and this is the honest part — most clinicians cannot tell a self-hosted Gemma from a BAA-covered Vertex Gemini at the demo. Heidi already runs in-region and signs BAAs. The deployment posture is a feature for procurement, not a moat for the clinician. It buys NHS-ready <i>language</i> without buying NHS-ready <i>certification</i> (see the integration roadmap).
5	Carryover-aware parent-facing artefacts (tone slider, reading level, mandatory AI disclosure footer)	Medium	18 months	The parent-facing surface is the genuinely under-served space — most competitors stop at the clinician note. Becomes a stronger moat if Sona invests in the parent portal experience.
6	Solo-practitioner-first product shape	Low	6 months	Pleasant, not defensible. A larger competitor can ship a "solo" SKU in a quarter. The real solo-practitioner moat is price + onboarding speed , not architecture.
7	GenUI / A2UI client architecture	Low	12 months	Important engineering choice, not a customer-visible moat. Do not list this in a pitch deck.
8	"Clinician-in-the-loop, always" principle	Low (as a moat)	0 months	Every competitor in 2026 says this. Calling it a moat dilutes the real moats above. It is a price of entry , not a differentiator.

The net: **moats 1, 2, and 3 are real and reinforce each other** — the loop, the posture, and the specialty-aware shape. Moat 4 deserves to be marketed but is a *trust* play, not a *defensibility* play. Moats 6, 7, and 8 should be dropped from the moat conversation; they belong in the principles / engineering footnotes.

3. The seams to add so Sona generalises without forking

Each seam is a single named change to a single subsystem. T-shirt sizes are relative to a 2-person team.

Seam 1 — Intake template engine (S → M)

Today: the 8 intake steps are hard-coded in the Flutter app; the API persists whatever the form sends. **Tomorrow:** a versioned, branded, conditionally-branching intake schema authored per practice and per specialty, stored as data, rendered by a single generic Flutter component. **What changes:** a new intake-template subsystem — a database table to hold the templates, an API service to fetch the right template for a given tenant, and a generic Flutter renderer that replaces today's hard-coded steps. **Why now:** it is the single biggest single-customer-tailoring risk and the prerequisite for every other generalisation. **Effort:** S to ship the engine with the existing 8-step paediatric template baked in as the first row; M to author a second template (adult voice) end-to-end. **Decision the founders own:** how schemas are *authored* — engineer-touch-only in source, a CMS-style editor (the design partner can edit their own), or a hybrid. Pick a side before the engineer starts.

Seam 2 — Plan and summary template library (M)

Today: one output-schema family for session plans, one tone-slider for summaries. **Tomorrow:** a *domain × specialty* template family — paediatric speech-sounds, adult voice, AAC introduction, feeding, stutter, social communication — each with its own prompt fragments, synthetic few-shot examples, and acceptance metrics. Specialty becomes the **key**, not a field. **What changes:** the session-plan and parent-summary services gain a prompt-template subsystem keyed by specialty. **Why now:** this is what stops the second SLT meeting from being "that's a kid's product".

Seam 3 — Vocabulary and statutory-framework pack (S)

Today: "Parent", "consult", "EHCP", "HCPC/RCSLT" are baked into strings throughout the code. **Tomorrow:** a per-tenant pack with three axes — recipient noun (parent / carer / client / patient), encounter noun (consult / appointment / session), statutory wrapper (EHCP for England, IDP for Wales, CSP for Scotland, Statement for NI, IEP for the US, NDIS for Australia, none), and regulator block (the UK pair HCPC+RCSLT vs ASHA in the US, SPA in Australia, SAC in Canada). One data file per pack; **no code** changes when adding a new jurisdiction. The Flutter app already has the localisation framework in place. **Why now:** trivial effort, large adoption-blocker removed. Cleanest "S" item on this list.

Seam 4 — Multi-jurisdiction plumbing extensions (M)

Today: an architecture decision already splits the UK and US data planes. **Tomorrow** for a real US/AU launch: payments (Stripe for cards vs GoCardless for UK direct debit, chosen per tenant), HIPAA business-associate agreements with US subprocessors (already in place for the AI-inference vendor, not yet for payment or email vendors), US-specific consent wording (the data model has a slot for it; needs to be populated per jurisdiction), and a small mapping layer so US fields line up with the US healthcare-data standard (USCDI) rather than the UK one (FHIR UK Core). **Why later:** Year 2+. The MVP and the design-partner pilot do not need this; build the seams so it is not a rewrite when it lands.

Seam 5 — Multi-tenant clinic mode (M → L)

Today: the team's working agreements in the MVP brief explicitly say *"Solo practitioner first. If a feature needs a 'team' concept to make sense, it's too early."* **Tomorrow:** a small-group-practice mode with **shared client list, per-clinician audit identity, shared template library, per-clinician inference budgets**. **When it becomes unavoidable:** the second time the design partner says "my associate is starting next month". Plan for it; do not build it yet. **Effort:** medium if the authentication and tenant model are set up correctly now (they appear to be — intake submissions are already tenant-scoped); large if a multi-clinician concept has to be retrofitted later.

Seam 6 — White-label / per-practice branding (S)

Today: Sona brand on every parent-facing artefact. **Tomorrow:** per-tenant logo, primary colour token, practice name in the form header, sender display name. The team's portal-first parent-comms model is unaffected. **Effort:** small — Flutter theming, a per-tenant branding store on the data side, and a logo-upload endpoint. **Why now:** it removes the "that's Speech Sanctuary's form" objection in one sprint.

Seam 7 — Specialty-aware persona library (S, ongoing)

Not a code seam, but a product seam. Today: four UK paediatric personas. Add three minimum before the first non-paediatric pitch: **adult voice (38yo gender-affirming voice), AAC introduction (5yo, non-verbal), feeding (3yo with progressed sensory profile)**. Personas drive the demo, the seed scripts, and the few-shot examples for the AI prompts. The cost of *not* adding them is showing up to the next SLT with a paediatric demo and learning nothing.

4. External-systems integration roadmap

Integrations are graded by **what they unlock** vs **what they cost**. "Cost" includes technical work, certification, and commercial / contracting overhead. Phasing is realistic for a two-person team plus one design partner.

UK (priority)

Integration	What it unlocks	Cost	Phase
Calendar (Google / Outlook / iCloud) — round-trip booking for the 20-min free consult	Closes the biggest "two products, two diaries" pain raised in the design-partner interview.	Low. Standard OAuth flows. Both Google Calendar API and Microsoft Graph have well-documented patterns.	MVP / Year 1
Video (Zoom, Whereby, Google Meet) deep-link generation — auto-attach a join URL to the consult slot	Removes the manual "send a Zoom link" step.	Low. Deep-link templating; no SDK needed.	MVP / Year 1
E-sign (DocuSign, Dropbox Sign) — T&Cs and consent on intake submit	Replaces the manual T&Cs PDF path; better-than-tickbox audit trail.	Medium. DocuSign for enterprise audit, Dropbox Sign cheaper. Both have free tiers below ~5 envelopes/mo.	Year 1
Payments (Stripe, GoCardless) — invoicing for the first session after triage	Direct revenue path. Stripe is faster to integrate; GoCardless is cheaper for UK direct debit recurring.	Medium. Dual-rail (Stripe for cards, GoCardless for DD) is what Cliniko already does.	Year 1
Cliniko / Jane App / Power Diary push — "Sona for Cliniko" — write Sona's intake + plan + summary into the existing PMS record	Distribution wedge: stop being "a Cliniko replacement" and become "the thing that makes your Cliniko worth keeping". This is the single highest-leverage GTM bet in this memo.	Medium. Cliniko has a public API; Jane has documented chart-attachment endpoints. PatientNotes already demonstrates that this distribution model works for AI tooling on top of Cliniko.	Year 1
NHS Login (parent identity vs magic link)	Parent-side identity matched to a real NHS account. Removes some magic-link friction.	Medium. Documented OIDC integration; vendor approval required but not gated on DSPT.	Year 2
NHS DSPT — Standards Met	Mandatory floor for anything NHS-adjacent. Confirmed Category 3 small-independent-provider self-assessment, no audit required, deadline 30 June 2026 .	Achievable for a 2-person company but requires real preparation — policy docs, asset register, Caldicott-Guardian-equivalent, backup evidence. Budget weeks not months.	Year 2 (mandatory prereq for everything below)

Integration	What it unlocks	Cost	Phase
GP Connect — Access Record: Structured (HL7 FHIR UK Core, NHS Digital)	Pull structured GP record into the intake brief — replaces the "type your GP's address" field. The single most valuable data enrichment available.	High. Prereqs: DSPT Standards Met + HSCN access + PDS compliance + IG model compliance + use-case approval + SCAL evidence + RBAC + a Clinical Safety Officer compliant with DCB0129 and DCB0160. Use-case approval is the gate before development starts; testing is a multi-Gate process.	Year 2+
NHS e-Referral Service (e-RS)	Only relevant if Sona ever takes NHS-funded private referrals.	High. Separate procurement path; not on the critical path for a private-practice product.	Year 2+ if positioning shifts to NHS-adjacent
NHS commissioning routes — G-Cloud 14, Spark DPS, NHS Digital Marketplace, direct ICB award	Path to NHS revenue. G-Cloud listing is cheap and quick; the real cost is the buyer-side ICB procurement cycle.	Low listing cost; high pipeline cost (3–9 months typical ICB procurement).	Year 2+

US (Year 2+)

Integration	What it unlocks	Cost	Phase
FHIR R4 + USCDI as the lingua franca	Anything US-integration-adjacent	Low if seam 4 is done correctly	Year 2+
EHR connectors — Epic App Orchard, Cerner / Oracle Health CODE, Athena Marketplace, eClinicalWorks, ModMed for therapy clinics	The real US distribution path. Epic alone has a multi-month onboarding.	High — each marketplace is months.	Year 2+
HIPAA BAA with GCP + subprocessors	Lawful US PHI processing	Partly done — the Google Cloud BAA is in place. The team has already concluded no transactional-email vendor will sign a BAA at this scale, so parent communications are delivered via an authenticated portal (the Epic MyChart model) rather than email.	Year 2+
IEP / IDEA workflow if school-district market is in scope	A different product really — schools-based SLT is enormous in the US.	High. Probably a Y3 conversation, not a Y2 one.	Year 3+

AU / CA (Year 2+, brief)

Integration	What it unlocks	Cost	Phase
My Health Record + FHIR AU Base; SPA accreditation	AU private SLT market	Medium — FHIR AU Base is well-documented; SPA is the regulator analogue.	Year 2+

Integration	What it unlocks	Cost	Phase
CA — SAC, provincial PHIPA / PIPEDA	Canadian SLT market	Medium — Jane App's own footprint shows the path.	Year 2+

Phased summary

Phase	UK	US	AU / CA
MVP (now)	Magic link, GCS exports, in-tenant audit	—	—
Year 1	Calendar + video + e-sign + payments + "Sona for Cliniko"	—	—
Year 2	NHS Login + DSPT Standards Met	FHIR R4 / USCDI scaffolding	—
Year 2+	GP Connect + e-RS + G-Cloud listing	EHR connectors (Epic, Cerner, Athena)	MHR / FHIR AU; SPA; PHIPA
Year 3+	ICB direct award	IEP / IDEA school-district workflows	—

5. Pricing options

Cost-side reality check

Per the team's architecture, Sona runs a single mid-tier GPU in the UK Google Cloud region for the AI inference. Each case triggers up to three inference passes: prep brief, session plan, parent summary. Per-case token budgets total roughly **10–15K tokens of input + 3K of output**, twice over (the parent summary runs a second simplification pass to hit a target reading age). A single L4 supports ~4 concurrent jobs; on-demand GCP L4 reservation rates put per-case inference cost in the **single-digit pence range at scale, mid-double-digit pence at MVP utilisation**. This is a directional figure — the founders should price-validate it against the bill before pricing materially below £49 / clinician / mo.

At the design partner's stated 3–8 cases / week / clinician, inference cost per clinician per month is in the **£10–£25 range, not the £200 range**. That keeps gross margin healthy at any of the per-clinician price points below.

Pricing options scored

Model	Target segment	Headline price	What's included	Gross margin (directional)	Pros	Cons
A — Per-clinician monthly subscription (recommended)	ASLTIP-member solo and 2–3 person practices	£79 / clinician / mo annual, £99 monthly	Unlimited intakes, AI drafting on a soft cap (e.g. 30 cases/mo), parent portal, audit trail	75–85% at typical volume	Predictable revenue; matches how Heidi, Jane, Clindoc, Smilenotes already bill; clinician understands the unit	Has to be "worth a Cliniko seat" — £79 is below Cliniko's own £45 + a £40 scribe combination, so it stays defensible
B — Per-case (intake → first-session loop)	Practices with very variable caseload (new graduates, returners)	£15–£25 / case	Single completed loop	Variable; lower at low volume	Aligns price to value; low-risk for new practices	Disincentivises high-volume usage; revenue too volatile to plan against; Sona's whole reason-to-exist is to <i>increase</i> throughput
C — Freemium	Top-of-funnel awareness; ASLTIP new-grads	First 5 cases / mo free; £49 / mo after	Free tier capped at 5 cases	Marketing line; not a revenue line at scale	Distribution; brings clinicians in via a friend	Cannibalises (a) — small practices stay in free tier; high free-tier inference cost

Model	Target segment	Headline price	What's included	Gross margin (directional)	Pros	Cons
D — NHS framework procurement (G-Cloud 14 / DPS)	NHS-adjacent later	£15–£25k / ICB / year (typical small-tool framework)	Whatever scope an ICB tenders	Higher gross margin but longer DSO	Bigger ACV; multi-clinician deployment	Different sales motion; gated on DSPT + GP Connect (see roadmap); is a Y2+ conversation
E — Per-practice flat + seat add-ons	The growth path for design partners onto small clinic mode	£149 / practice base + £49 / additional seat	Base = 1 clinician's loops; seats = additional clinicians	Highest gross margin at clinic scale	Supports the multi-tenant clinic-mode seam; good investor narrative on expansion revenue	More complex billing; needs the clinic-mode work to land first

Recommended starting model

Pick A with a tiered headline: £39 "Starter" (cap 5 cases/mo; new-grad / returner / ASLTIP-affiliate path), **£79 "Practice"** (the default; cap 30 cases/mo), £149 "Practice+" (no cap; priority support). Move to model E (per-practice + seats) **only when the second clinician at the design-partner practice signs**, not before. Avoid (B) and (C) until the price-sensitivity of the ASLTIP base is measured. Pursue (D) only after DSPT in Year 2.

Why £79 and not £49 (Clindoc) or £99 (Heidi Pro era): Sona is not a scribe-only product, so it should not be priced like one. £79 sits exactly between Smilenotes (£5, no AI) and Heidi Clinician (~£120 GBP, scribe-only) and quantifies the bundle: PMS-light + AI loop + UK residency + audit. The price-anchor effect of being deliberately "between" two known reference points helps the discovery call.

Recommendations — what to do in the next 30 / 60 / 90 days

Ordered by leverage. Day windows are windows, not deadlines — the slice that fits the team's actual capacity is the right one.

1. **(Days 1–14) Run the one-therapist-fit audit with a second SLT.** Specifically: an SLT outside the design-partner practice, ideally adult voice or AAC, ideally an ASLTIP member. Use the audit table above as the agenda; ask them to walk the demo and call out leaks. **Owner:** founders. **Effort:** 1 working day of prep + 90 minutes interview. **Metric:** at least 5 leaks confirmed, at least 2 contradicted (i.e. *not* a problem for that SLT).
2. **(Days 1–30) Ship seam 6 (white-label branding) and seam 3 (vocabulary pack).** The smallest possible product changes that remove the "that's Speech Sanctuary's form" objection. **Owner:** engineering. **Effort:** S each. **Metric:** the demo runs end-to-end under a fictional second-practice name, with no Speech Sanctuary string visible.
3. **(Days 15–60) Pick one of: intake template engine (seam 1) OR the AI-loop completion — i.e. shipping the prep brief, session plan, and parent summary end-to-end with real model output rather than stub content. Not both.** This is the single biggest engineering decision in the next quarter. The template engine generalises the product. The AI loop completes the demo's most fragile claims. The wrong call is to try to do both at half-speed. **Owner:** founders. **Effort:** decision-only this sprint; M–L next quarter.
4. **(Days 30–60) Stand up a costed DSPT readiness plan with a specialist consultancy.** Three quotes; choose one; book the work for the autumn. DSPT is the gate for everything NHS-adjacent and is achievable as a self-assessment — but real evidence files, an asset register, and a backup-restore drill take preparation. **Owner:** founders. **Effort:** half a day of vendor calls. **Metric:** chosen consultancy, scoped quote.
5. **(Days 30–90) Approach ASLTIP about a member-benefit slot or directory listing.** ASLTIP markets to 1,800 of the highest-quality UK private SLTs and explicitly hosts events that surface tooling (OEPR, the most clinically aware UK competitor, markets via ASLTIP events). A discounted listing or a co-marketed launch is the single highest-leverage GTM move available. **Owner:** founders. **Effort:** S. **Metric:** a meeting with the ASLTIP board.
6. **(Days 60–90) Build "Sona for Cliniko" — write the intake brief, plan, and summary into a Cliniko patient record.** Stop competing with Cliniko; complement it. PatientNotes' existence on Cliniko's connected-apps marketplace is proof the distribution model works. This converts every Cliniko-using ASLTIP member into a prospect rather than a defended account. **Owner:** engineering. **Effort:** M. **Metric:** one paying clinician using both products in production.

Open questions for the founders

In priority order; each blocks something downstream.

1. **Is the next bet generalisation (intake template engine) or completion (AI loop end-to-end)?** Both at half-speed is the failure mode (Recommendation 3).
2. **What is the actual paying-customer hypothesis: paediatric-only private SLT, or all-of-SLT-on-launch?** The honest answer changes everything in section 1 and the pitch deck. The current personas and the EHCP wiring suggest paediatric-only is the truthful framing.
3. **Are you willing to be priced below Heidi and above Smilenotes, or do you want a wedge price?** £79 / clinician / mo is a defensible middle. A £29 wedge price changes the product (closer to Smilenotes) and the moat conversation (less "premium UK posture", more "low-cost ASLTIP default").
4. **DSPT — Year-2 commit or Year-3 "wait and see"?** A real DSPT commit unlocks GP Connect, NHS Login, and the G-Cloud listing in series. A wait-and-see lets the product mature on private-pay first. Both are defensible; the deck cannot be vague about which one.
5. **OEPR (Octopus EPR) — competitor, co-existent, or eventual partner?** OEPR is the most clinically aware UK SLT-specific competitor surfaced in this review. They market via ASLTIP. The interesting future is a Sona-on-OEPR integration, not a Sona-vs-OEPR fight. The founders need to decide whether to reach out now or later.
6. **Multi-tenant clinic mode — explicit non-goal for 2026, or a Q4 2026 build?** The design-partner relationship will demand it once a second clinician joins the design-partner practice.
7. **Persona expansion — agreed scope of three new personas (adult voice, AAC, feeding) before the next external pitch?** Without this, every non-paediatric meeting is uphill.

Sources

Internal source artefacts (non-public, team repository)

- **MVP brief** — original positioning, scope, design-partner interview synthesis.
- **Feedback-demo brief** — the "real today vs 90-day vision" split that underpins every honesty claim in this memo.
- **Intake-form specification** — the 8-page Speech Sanctuary parent questionnaire the MVP form is built from.
- **AI design notes for triage, session-plan, and parent-summary** — output schemas, prompt structures, target metrics.
- **Five architecture decision records** covering data residency, the Flutter front-end, the self-hosted AI model, environment access controls, and portal-first parent communications.
- **Architecture overview** — UK/US split, region pinning, business-associate-agreement chain.
- **Four synthetic demo personas** — 4-year-old speech-sounds, 7-year-old stutter with EHCP, 11-year-old social communication, 3-year-old feeding.
- **Flutter app and API source for the parent intake form** — the current hard-coded eight-step flow.

Web sources (URL + access date, all 2026-05-28)

- Heidi Health Clinician plan (\$150/user/mo) — veroscribe.com/blog/heidi-health-review-2026
- Heidi Health, second pricing source — deepcura.com/resources/heidi-health-review
- Heidi Health SLP assessment template — heidihealth.com/templates/review-speech-pathology-assessment
- Jane App AI Scribe (\$15/mo, US+CA only) — jane.app/features/charting-ai-scribe
- Jane App AI Scribe — guide — jane.app/guide/ai-scribe
- Cliniko connected-apps and AI marketplace — [wisevu.com Cliniko Review](https://wisevu.com/Cliniko%20Review)
- Cliniko form upload — [help.cliniko.com upload-form-templates](https://help.cliniko.com/upload-form-templates)
- Clindoc UK pricing (£40 / £5 active day) — clindoc.ai/pricing
- PatientNotes for speech therapists — patientnotes.app/professions/speech-therapist-st
- Octopus EPR (OEPR) — oepr.co.uk
- Smilenotes UK — smilenotes.co.uk/profession/speech-therapy-software
- ASLTIP membership statistics — asltip.com and asltip.com/why-join
- NHS GP Connect Access Record: Structured (FHIR API) — digital.nhs.uk/developer/api-catalogue/gp-connect-access-record-structured-fhir
- NHS GP Connect — Consumer Assurance Process — github.com/nhsconnect/gpc-consumer-support
- NHS DSPT — Standards Directory — standards.nhs.uk/published-standards/data-security-and-protection-toolkit
- DSPT 2025/26 small-org guide — dsptready.co.uk/blog/dspt-complete-guide
- Evalian DSPT 2025 guide — evalian.co.uk/the-data-security-and-protection-toolkit-dspt
- Chatter Labs — adjacent SLT AI landscape — chatter-labs.com/blog/slt-private-practice-setup-tips/ai-speech-therapy

Engineering footnotes

- **Intake template engine concrete shape.** A new database table holds versioned, per-tenant, per-specialty intake schemas (id, tenant, specialty, locale, version, steps, branches, vocabulary-pack reference). The Flutter renderer reads the steps and dispatches to widget types keyed by field kind (text, date, multi-select, yes/no-with-detail, composite GP details, consent block). The current 8-step form ships as the first row, version `v1-paediatric-uk`. The AI context window for prep brief / session plan / summary already operates on the form's answers blob rather than on hard-coded step IDs — so the template engine is purely a client + per-tenant authoring change, not an AI-context change.
- **Self-hosted Gemma economics.** L4 24 GB INT4 AWQ Gemma 3 27B in the London Google Cloud region. Per the AI architecture decision, 4 concurrent jobs / GPU; per-job P95 ≤ 8 s (prep) + ≤ 15 s (plan) + ≤ 20 s (summary). At an L4 reservation rate and pilot caseload (design partner ≈ 8 cases/wk ≈ 32 /mo), per-clinician inference cost is in the **£10–£25/mo** range, dominated by reservation rather than utilisation. At 50 clinicians on the same L4, marginal per-clinician cost compresses below £5.
- **GP Connect — concrete onboarding shape.** Per NHS Digital, retrieved 2026-05-28: prereqs are DSPT Standards Met + HSCN access + PDS-compliant identity (or third-party PDS proxy) + IG model compliance + RBAC + CSO appointment compliant with DCB0129 + DCB0160. Onboarding = use-case submission $\rightarrow$ approval $\rightarrow$ 6-month dev window $\rightarrow$ SCAL evidence $\rightarrow$ multi-Gate testing (test environments per Gate) $\rightarrow$ go-live. The self-hosted AI stack is irrelevant to GP Connect onboarding; the API and tenant-routing layers are where the FHIR R3 client lives.
- **OEPR and Cliniko on the same diagram.** OEPR is a domain-specific record system; Cliniko is a generalist PMS with an AI-scribe marketplace. Sona's natural position is on the **A2UI / read-modify side**, not on the records side — i.e. Sona writes intake/plan/summary objects into whichever record system the practice runs (OEPR via TOMs-aware mapping; Cliniko via patient-form templates + chart notes). Treat OEPR as a future read/write target, not a competitor to be displaced.